

PART 2 — JUDGE Q&A MASTER PREPARATION

Now comes the **more important part**.

Your PPT can be excellent, but a team can lose an internal round through Q&A.

I'm going to prepare you for questions beyond the 12 you already have.

CATEGORY A — PROBLEM QUESTIONS

Q1. "Exactly what problem are you solving?"

Answer

"The problem is not lack of MPLADS data. The problem is converting large volumes of existing administrative and project evidence into scalable, explainable intelligence that helps authorities identify which works need attention first."

If they want one line:

"Data exists; scalable intelligence for prioritization is the gap we are addressing."

Q2. "Why can't officials just look at the existing eSAKSHI dashboard?"

★ VERY LIKELY

"eSAKSHI is the underlying digital workflow and data ecosystem. We are not claiming that it cannot display information. Our focus is different: instead of expecting an official to manually identify complex patterns across multiple

dimensions, our intelligence layer automatically generates cross-signal risk indicators, prioritizes works and explains why they deserve attention."

Then:

"eSAKSHI manages the workflow; our proposed layer focuses on intelligence and prioritization."

That distinction is extremely important.

Q3. "Isn't this already available in eSAKSHI?"

Answer

"Some individual information and public-facing features already exist. Our novelty is not claiming that MPLADS has no monitoring or feedback mechanism. Our contribution is the **cross-signal intelligence layer** that combines compliance, financial, operational, relationship and evidence signals into an explainable risk profile with evidence strength and data coverage."

Q4. "Who exactly will use your system?"

Answer

"Primary users would be authorized monitoring and administrative authorities such as the Ministry, State Nodal Authorities and District Authorities. The system is decision support; it doesn't replace their authority."

Your PPT also identifies authorized investigators, citizens and genuine implementers as other beneficiaries.

CATEGORY B — AI QUESTIONS

Q5. "Where exactly is the AI in your project?"

Don't say "everywhere."

Say:

"We use a hybrid intelligence architecture. Deterministic rules handle compliance and data-quality conditions. Statistical methods handle peer benchmarking. Isolation Forest and DBSCAN can identify multidimensional anomalies. OpenCV and perceptual hashing support evidence analysis, while NetworkX handles relationship analytics. So AI/ML is used specifically where it adds value."

Q6. "Why Isolation Forest?"

Answer

"Because we don't have reliable confirmed-fraud labels for the prototype. Isolation Forest is suitable for unsupervised anomaly detection because it identifies observations that are unusual compared with the overall data distribution."

Q7. "Why DBSCAN?"

Answer

"DBSCAN is useful for finding clusters and outliers without requiring us to specify the number of clusters beforehand. In our context it can support anomaly analysis of patterns such as coordinated feedback behaviour."

Important:

Don't claim DBSCAN is your fraud detector.

It's a **supporting anomaly-analysis method**.

Q8. "Why not Random Forest?"

Answer

"Random Forest is primarily a supervised learning approach. To train it properly for fraud classification, we would need reliable labelled examples of fraud and non-fraud. We don't want to manufacture those labels. For our current setting, unsupervised anomaly detection is more defensible."

🔥 Excellent answer.

Q9. "Why not use deep learning?"

Answer

"Deep learning is not automatically better. Our primary requirements are explainability, limited labelled data and reliable MVP demonstration. We therefore use simpler techniques where appropriate. More advanced models such as GraphSAGE are part of our future research once sufficient validated data becomes available."

Q10. "How exactly does your risk score work?"

Answer

"Each dimension receives a score from 0 to 100. Our frozen weighting is Financial 25%, Compliance 20%, Operational 20%, Relationship 15% and Evidence 20%. The overall priority is a transparent weighted combination, with documented renormalization if a dimension is unavailable."

Then, if needed:

$$\text{Risk} = 0.25F + 0.20C + 0.20O + 0.15R + 0.20E$$

Very important

If your prototype implements a different formula, **change the prototype to match this formula before demo day.**

Never let PPT mathematics and software mathematics disagree.

CATEGORY C — "FRAUD"

QUESTIONS

Q11. "Can your system detect fraud?"

THE MOST IMPORTANT ANSWER

"It detects potential risk indicators, not confirmed fraud."

Then:

"A single anomaly is not treated as fraud. Multiple independent signals can increase verification priority. The final determination remains with authorized human authorities."

Q12. "What if a project is genuinely expensive?"

"That's exactly why we use peer benchmarking as an anomaly signal rather than a fraud conclusion. A higher-than-peer cost could have legitimate explanations such as project scale, geography or technical requirements. The system therefore recommends verification rather than accusing the project."

Q13. "What if payment is ahead of progress?"

"It becomes a financial or operational risk signal. There may be legitimate reasons depending on project stage and payment structure. That's why the system doesn't automatically call it fraud."

Q14. "What if your model falsely flags an honest contractor?"

"The output is a verification priority, not an accusation. The officer reviews the underlying evidence and can record the outcome as valid concern, false positive or requiring further investigation."

Q15. "Can your system predict that an MP will commit fraud?"

Answer:

"No, and deliberately not. We do not predict individual criminal intent. We predict project-level implementation risks such as delay, cost escalation and compliance slippage."

This answer should be **instant**.

CATEGORY D — DATA QUESTIONS

Q16. "Where did you get the dataset?"

Answer

"For the hackathon prototype, we are using representative synthetic MPLADS-style data because field-level production eSAKSHI access has not been established for us."

Then immediately:

"We have structured it around realistic entities such as works, sanctions, payments, progress, agencies, vendors and evidence."

Q17. "Then how can you claim your system works?"

Answer

"We don't claim real-world fraud detection accuracy from synthetic data. We demonstrate functional detection using controlled test cases where known anomalies are deliberately planted."

Killer line:

"Synthetic testing proves prototype functionality, not real-world fraud-detection accuracy."

Q18. "What happens when real data comes?"

"The synthetic dataset can be replaced with authorized production data because the architecture separates ingestion, preprocessing, feature engineering and intelligence modules. We don't need to redesign the entire system."

Q19. "What if some data is missing?"

★ **Strong answer**

"Our architecture is data-availability-agnostic. A missing field doesn't automatically become low risk. We show **Data Coverage** separately and activate only the intelligence modules supported by available data."

Q20. "Why Data Coverage?"

"Because absence of evidence shouldn't accidentally become evidence of safety. If relationship or GPS information isn't available, the user should know that the assessment was made without it."

🔥 This is a **very strong conceptual answer**.

CATEGORY E — TECHNICAL IMPLEMENTATION

Q21. "How do you detect duplicate images?"

"We use perceptual hashing. Instead of comparing only the exact binary file, the image is transformed into a perceptual representation, allowing us to identify visually similar or reused images even when basic file properties change."

If asked about library:

"ImageHash, with OpenCV where additional image processing is required."

Q22. "What does OpenCV actually do?"

"It supports image preprocessing and computer-vision analysis. In our evidence module it can help normalize and inspect images before similarity or consistency analysis."

Don't claim OpenCV magically detects fraud.

Q23. "How does peer benchmarking work?"

Answer

"A work is compared with comparable works based on available authorized attributes such as category, geography, sanctioned amount and scale or quantity where available. We calculate deviation from the peer baseline."

Example:

"If comparable works have a median cost of ₹10 lakh and the current work is ₹16 lakh, the +60% deviation becomes a financial anomaly signal."

Q24. "Why not simply use the average?"

"The mean can be strongly affected by extreme values. For peer cost benchmarking, median and robust statistics can provide a more stable baseline. ML outlier detection can then act as an additional signal."

Excellent.

Q25. "How does relationship intelligence work?"

"We construct a graph where entities such as MP, constituency, agency, vendor, work and payment can become nodes, with relationships such as recommended, executed-by, paid-to or located-in represented as edges. Network analysis can then identify concentration or repeated relationship patterns."

Then:

"We only use fields that are actually authorized and available."

Q26. "Why NetworkX?"

"For an MVP, NetworkX is lightweight, transparent and sufficient for graph construction and basic network analytics. A learned graph model such as GraphSAGE becomes more appropriate as future research when sufficient validated network data exists."

CATEGORY F — EVIDENCE FUSION

Q27. "What exactly is Evidence Fusion?"

Answer

"Evidence Fusion is the stage where outputs from different intelligence modules are combined rather than considered independently."

Example:

Payment mismatch
+
Cost anomaly
+
Delay
+
Evidence inconsistency
↓
Evidence Fusion
↓
Higher verification priority

Then:

"The important point is that corroboration increases priority; it does not automatically establish fraud."

Q28. "How do you know two signals are independent?"

This is a **harder judge question**.

Answer:

"We distinguish signals by their evidence source or analytical dimension. For example, payment-progress mismatch is financial/operational, while image reuse is evidence-based. In the prototype we treat cross-category corroboration as stronger than repeated variations of the same underlying signal. In production, independence and reliability weights would need to be calibrated using validated historical outcomes."

Very mature answer.

Q29. "How do you calculate Evidence Strength?"

"It is not simply the number of alerts. It considers the number of signals, their independence, and the reliability of the underlying evidence. A single weak signal may be an observation; multiple independent credible signals can justify higher verification priority."

CATEGORY G — CITIZEN VERIFICATION

Q30. "Why do you need citizens? Isn't this a government system?"

"Citizen verification is optional, not load-bearing. Administrative records remain the primary source. The QR mechanism provides an additional ground-level signal when available."

Q31. "What if people abuse the feedback?"

"We don't treat individual citizen submissions as truth. We can analyze abnormal submission rates or coordinated patterns using anomaly detection, and the citizen signal remains only one input to evidence fusion."

Q32. "Why don't you require OTP?"

"For the MVP we want a minimal-friction verification flow. We deliberately avoid making account creation or OTP a dependency. Production identity, privacy and retention requirements would be governed according to the applicable institutional and legal requirements."

CATEGORY H — FEASIBILITY

Q33. "Can this actually be deployed by the government?"

Don't say:

"Yes, immediately."

Say:

"The architecture is technically deployable, but production deployment would require authorized data access, government-approved infrastructure, security review, privacy and governance controls, and validation with domain experts. Our current deliverable is a prototype demonstrating the intelligence architecture."

Q34. "How much will it cost?"

"Our prototype target is ₹0 using open-source software and local/free-tier infrastructure. Production cost cannot responsibly be estimated until data volume, infrastructure, security and government deployment requirements are defined."

Q35. "Can it handle millions of records?"

"The architecture is designed to scale from the prototype database toward PostgreSQL and government-approved infrastructure. The analytical modules can be processed in batches or distributed as data volume grows. We have intentionally separated ingestion, processing and intelligence layers."

CATEGORY I — LIMITATIONS

Q36. "What's your biggest limitation?"

Best answer:

"The biggest current limitation is validated real-world data."

Then:

"Because we don't have established production field-level access, our prototype uses synthetic data. Therefore our current results demonstrate functionality rather than real-world fraud-detection accuracy."

Don't try to hide this.

Q37. "What is your second biggest limitation?"

"False positives. An unusual project doesn't necessarily indicate wrongdoing. That's why we use multi-signal corroboration, evidence strength, data coverage and human review."

Q38. "What if the system misses fraud?"

"No anomaly-detection system can guarantee zero false negatives. Our architecture reduces dependence on a single detector by combining different evidence dimensions. But production evaluation would require historical validated cases and continuous monitoring."

CATEGORY J — COMPETITION / USP

Q39. "Why should we select you?"

Memorize this.

"Because we are not simply building another fraud dashboard or attaching an ML model to MPLADS data.

Our core contribution is an explainable multi-signal intelligence layer.

We combine compliance, financial and operational behaviour, relationships and available evidence.

Then we don't just output a risk score. We show:

What needs attention, why it needs attention, how strong the evidence is, how much data was available, and what action should be considered.

And the final decision remains with the authorized human authority."

Q40. "What is your USP in one sentence?"

 **Best answer:**

"Our USP is evidence-backed risk prioritization rather than black-box fraud prediction."

Q41. "What's the difference between your system and a normal dashboard?"

"A dashboard primarily presents information. Our platform analyzes information, correlates signals, prioritizes works, explains the reasons and recommends the next human action."

Q42. "What's the difference between your system and a normal anomaly detector?"

"A conventional anomaly detector may identify one unusual observation. Our system combines multiple intelligence dimensions and converts those signals into an explainable project-level risk profile with evidence strength, data coverage and recommended action."

PART 3 — THE 10 QUESTIONS I EXPECT MOST TOMORROW

If you have **very little time tonight**, memorize these.

 1.

What problem are you solving?

Data exists. The challenge is scalable intelligence for prioritizing which works need attention first.

 2.

What's innovative?

Multi-source evidence fusion with explainable risk prioritization.

 3.

Are you detecting fraud?

No. We detect risk indicators and prioritize human verification.

 4.

Where is the AI?

Isolation Forest/DBSCAN for anomaly analysis, CV for evidence, graph analytics for relationships, while deterministic rules handle deterministic compliance checks.

 5.

Where is your data?

Representative synthetic data because production field-level eSAKSHI access hasn't been established.

 6.

How did you test it?

Controlled planted anomaly cases.

 7.

How do you handle false positives?

Multi-signal corroboration + evidence strength + human review + outcome feedback.

🔥 8.

Why not just eSAKSHI?

eSAKSHI is the digital workflow/data ecosystem; we're proposing an intelligence layer focused on cross-signal analysis and prioritization.

🔥 9.

Why should we select you?

We turn existing administrative data into explainable evidence-backed prioritization rather than simply building another dashboard.

🔥 10.

Explain your project in one sentence.

"We are an explainable intelligence layer over authorized MPLADS/eSAKSHI data that combines multiple risk and evidence signals to identify which development works should receive human verification priority."

PART 4 — QUESTIONS THAT CAN CATCH YOUR TEAM OFF-GUARD

These are the ones I would **actively rehearse**, because judges often use them to test whether the team actually built the project.

"Show me exactly what happens when one signal fires."

Answer:

"It initially becomes an observation or risk signal. It doesn't automatically become a high-priority case. Its contribution depends on the dimension score, reliability and corroboration with other available signals."

"What happens if there is only one data source?"

"The system can still operate using available modules, but Data Coverage will reflect the limitation. We don't pretend the assessment has the same confidence as a case supported by multiple independent evidence categories."

"Why 25% financial and 20% evidence?"

"These are our transparent prototype weights, chosen to reflect the importance of those dimensions while avoiding dominance by a single category. They are not claimed to be universally optimal. In production, weights should be validated and recalibrated using domain expertise and historical outcomes."

 **This is much better than pretending 25% is scientifically proven.**

"Who decided the threshold for HIGH risk?"

Say:

"For the prototype, thresholds are configurable rather than treated as universal truths. Production thresholds would need validation using historical cases, expert review and operational requirements."

"Can your model be manipulated?"

"Any data-driven system can be exposed to manipulated inputs. Our mitigation is cross-signal corroboration, data-quality checks, evidence comparison and human review rather than trusting a single field or model."

"What happens after an officer marks a false positive?"

"The outcome is recorded. In the MVP this supports auditability and analysis. In the future, verified outcomes can be used for threshold and model recalibration."